

# Practice Exam

Scoring:

=====

0-11: fail (1),  
12-14: pass (2),  
15-17: satisfactory (3),  
18-20: good (4),  
21-24: excellent (5).

The file [forestfires.txt](#) contains data about forest fires in Portugal. Write a program that estimates the value of the *area* feature from the other features! Apply 5-fold cross-validation! The error metric should be the root mean squared error (RMSE).

Subproblems:

- (a) Perform initial data analysis! Prepare the input matrix  $X$  and the target vector  $y$ ! [4 points]
- (b) Create a t-SNE based visualization of the input features! [4 points]
- (c) Introduce at least one new feature that is engineered from the original ones. [4 points]
- (d) Create a non-tree based estimator and optimize its hyper-parameters! [6 points]
- (e) Create a tree based estimator and optimize its hyper-parameters! [6 points]

**(a) Perform initial data analysis! Prepare the input matrix  $X$  and the target vector  $y$ ! [4 points]**

```
In [5]: # Read data from CSV file.  
import pandas as pd  
fname = 'forestfires.txt'  
df = pd.read_csv(fname)
```

```
In [6]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 517 entries, 0 to 516
Data columns (total 13 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    X      517 non-null     int64  
1    Y      517 non-null     int64  
2    month  517 non-null     object  
3    day    517 non-null     object  
4    FFMC   517 non-null     float64 
5    DMC    517 non-null     float64 
6    DC     517 non-null     float64 
7    ISI    517 non-null     float64 
8    temp   517 non-null     float64 
9    RH     517 non-null     int64  
10   wind   517 non-null     float64 
11   rain   517 non-null     float64 
12   area   517 non-null     float64 
dtypes: float64(8), int64(3), object(2)
memory usage: 52.6+ KB

```

- there are no null values
- month and day are strings

```

In [7]: # columns stats
df.describe().T

```

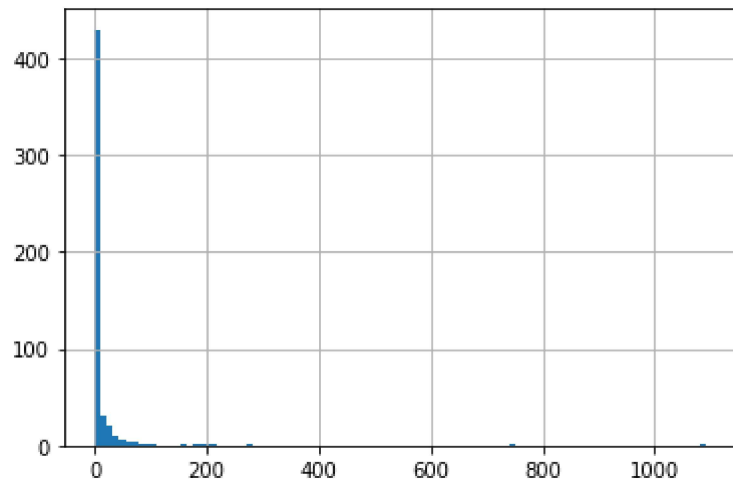
Out[7]:

	count	mean	std	min	25%	50%	75%	max
<b>X</b>	517.0	4.669246	2.313778	1.0	3.0	4.00	7.00	9.00
<b>Y</b>	517.0	4.299807	1.229900	2.0	4.0	4.00	5.00	9.00
<b>FFMC</b>	517.0	90.644681	5.520111	18.7	90.2	91.60	92.90	96.20
<b>DMC</b>	517.0	110.872340	64.046482	1.1	68.6	108.30	142.40	291.30
<b>DC</b>	517.0	547.940039	248.066192	7.9	437.7	664.20	713.90	860.60
<b>ISI</b>	517.0	9.021663	4.559477	0.0	6.5	8.40	10.80	56.10
<b>temp</b>	517.0	18.889168	5.806625	2.2	15.5	19.30	22.80	33.30
<b>RH</b>	517.0	44.288201	16.317469	15.0	33.0	42.00	53.00	100.00
<b>wind</b>	517.0	4.017602	1.791653	0.4	2.7	4.00	4.90	9.40
<b>rain</b>	517.0	0.021663	0.295959	0.0	0.0	0.00	0.00	6.40
<b>area</b>	517.0	12.847292	63.655818	0.0	0.0	0.52	6.57	1090.84

- standard deviations are quite different
- max(area) is much bigger than the 75-percentile.

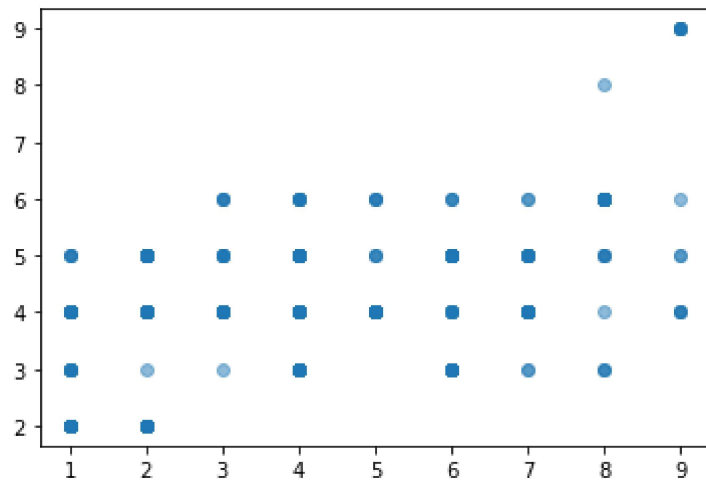
```
In [15]: # histogram of area
df['area'].hist(bins=100)
```

```
Out[15]: <AxesSubplot: >
```



```
In [19]: # scatter plot of (X, Y)
import matplotlib.pyplot as plt
plt.scatter(df['X'], df['Y'], alpha=0.5)
```

```
Out[19]: <matplotlib.collections.PathCollection at 0x7f3098439490>
```



```
In [21]: # possible values of month
df['month'].unique(), df['month'].nunique()
```

```
Out[21]: (array(['mar', 'oct', 'aug', 'sep', 'apr', 'jun', 'jul', 'feb', 'jan',
        'dec', 'may', 'nov'], dtype=object),
        12)
```

```
In [22]: # possible values of day
df['day'].unique(), df['day'].nunique()
```

```
Out[22]: (array(['fri', 'tue', 'sat', 'sun', 'mon', 'wed', 'thu'], dtype=object), 7)
```

```
In [28]: columns = ['X', 'Y', 'FFMC', 'DMC', 'DC', 'ISI', 'temp', 'RH', 'wind', 'rain']
X = df[columns].values # input matrix
y = df['area'].values # target vector
```

```
In [32]: X.shape, y.shape, X.mean(), y.mean()
```

```
Out[32]: ((517, 10), (517,), 83.46644100580271, 12.847292069632493)
```

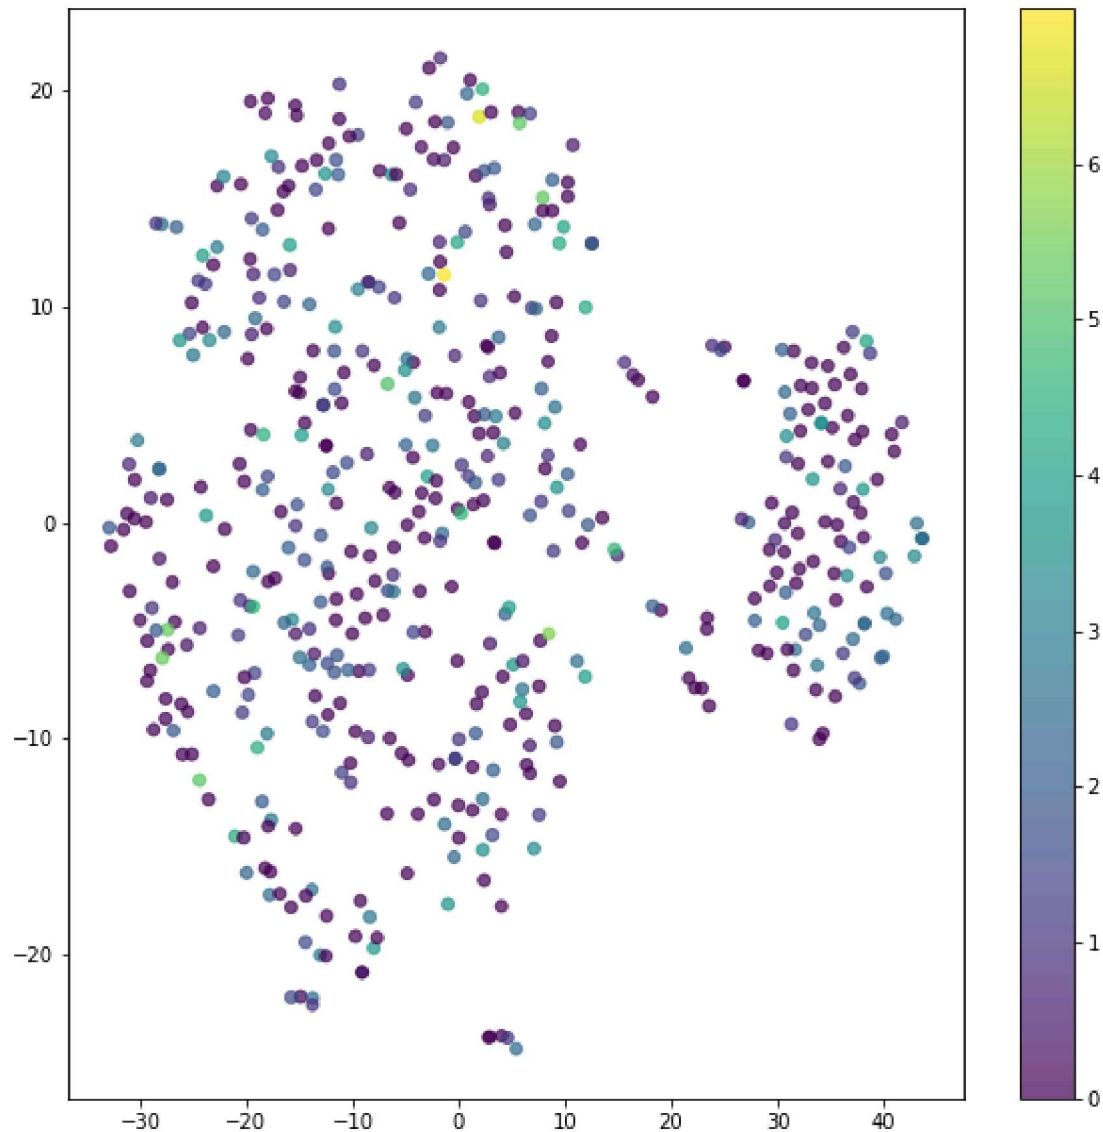
```
In [35]: from sklearn.preprocessing import StandardScaler
X2 = StandardScaler().fit_transform(X) # standardize input matrix
```

## (b) Create a t-SNE based visualization of the input features! [4 points]

```
In [39]: from sklearn.manifold import TSNE
Z = TSNE(random_state=42).fit_transform(X2)
```

```
In [51]: import numpy as np
plt.figure(figsize=(10, 10))
plt.scatter(Z[:, 0], Z[:, 1], c=np.log(y + 1), alpha=0.7)
plt.colorbar()
```

```
Out[51]: <matplotlib.colorbar.Colorbar at 0x7f3089c08580>
```



(c) Introduce at least one new feature that is engineered from the original ones. [4 points]

```
In [58]: # convert month name to month index
month_idx = {
    'jan': 1, 'feb': 2, 'mar': 3, 'apr': 4, 'may': 5, 'jun': 6,
    'jul': 7, 'aug': 8, 'sep': 9, 'oct': 10, 'nov': 11, 'dec': 12,
```

```
}
df['month_idx'] = [month_idx[m] for m in df['month']]
```

```
In [63]: # convert weekday name to weekday index
day_idx = {
    'mon': 1, 'tue': 2, 'wed': 3, 'thu': 4, 'fri': 5, 'sat': 6, 'sun': 7
}
df['day_idx'] = [day_idx[d] for d in df['day']]
```

### (d) Create a non-tree based estimator and optimize its hyper-parameters! [6 points]

```
In [78]: # Ridge regression with default parameters.
from sklearn.linear_model import Ridge
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error

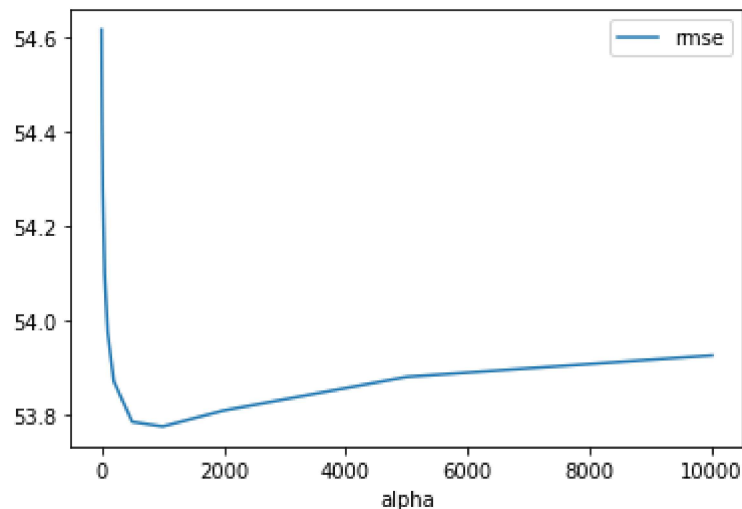
def evaluate(re, X, y):
    cv = KFold(shuffle=True, random_state=42)
    scores = []
    for tr, te in cv.split(X):
        re.fit(X[tr], y[tr])
        yhat = re.predict(X)
        rmse = mean_squared_error(y[te], yhat[te])**0.5
        scores.append(rmse)
    return np.mean(scores)

X3 = StandardScaler().fit_transform(df[columns + ['month_idx', 'day_idx']].values)
evaluate(Ridge(), X3, y)
```

Out[78]: 54.58729606435613

```
In [88]: # Hyperparameter optimization.
alphas = [0.01, 0.02, 0.05, 0.1, 0.2, 0.5, 1, 2, 5, 10, 20, 50, 100, 200, 500, 1000, 2000, 5000, 10000]
res = []
for alpha in alphas:
    res.append({
        'alpha': alpha,
        'rmse': evaluate(Ridge(alpha=alpha), X3, y)
    })
df_res = pd.DataFrame(res).set_index('alpha')
df_res.plot()
```

Out[88]: <AxesSubplot: xlabel='alpha'>



In [91]: `df_res['rmse'].idxmin(), df_res['rmse'].min()`

Out[91]: (1000.0, 53.77504142548006)

### (e) Create a tree based estimator and optimize its hyper-parameters! [6 points]

In [95]: `from sklearn.ensemble import GradientBoostingRegressor`

`evaluate(GradientBoostingRegressor(random_state=42), X3, y)`

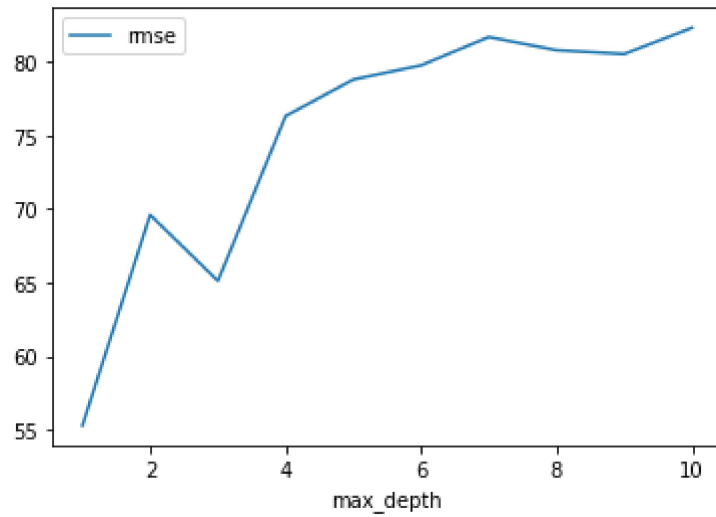
Out[95]: 65.06556168485034

In [96]: `# Hyperparameter optimization.`

```
res = []
for max_depth in range(1, 11):
    print(max_depth, end=' ')
    res.append({
        'max_depth': max_depth,
        'rmse': evaluate(GradientBoostingRegressor(max_depth=max_depth), X3, y)
    })
df_res = pd.DataFrame(res).set_index('max_depth')
df_res.plot()
```



Out[96]: <AxesSubplot: xlabel='max\_depth'>



In [97]: `df_res['rmse'].idxmin(), df_res['rmse'].min()`

Out[97]: (1, 55.321957821450226)